MULTIZ 11.4 scores each yama cell from column counts, not K × L row pairs

A 11.3: every row pair, in every cell

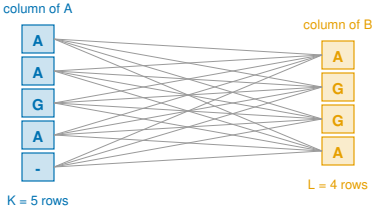

B 11.4: summarise each column once, then combine

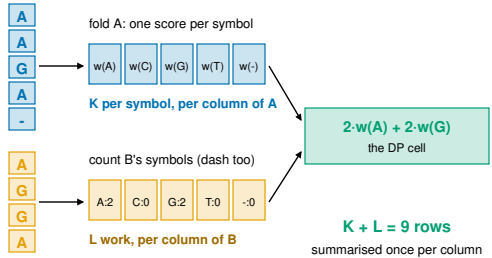

Gap sums are counted the same way. A row's gap cost depends only on whether it holds a dash in this column and in the previous one, so each row falls into one of 4 dash classes. All terms are integers: the totals equal the pair sums exactly.

C One deep merge, as recorded in multiz#28

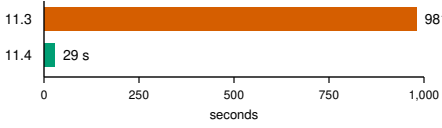

Blocks of about 35 rows against 43, cs10 included:

11.3: every pair of the rows yama aligns, in every cell

11.4: those rows summarised once per column

33.8× faster on this merge,

identical output (single upstream runs)